

# CS689: COMPUTATIONAL LINGUISTICS FOR INDIAN LANGUAGES TERM EMBEDDING

Arnab Bhattacharya  
`arnabb@cse.iitk.ac.in`

Computer Science and Engineering,  
Indian Institute of Technology, Kanpur  
<http://web.cse.iitk.ac.in/~cs689/>

2<sup>nd</sup> semester, 2023-24  
Tue 10:30–11:45, Thu 12:00–13:15 at RM101/KD102

# Embedding Models

- Models for *embedding* terms to vector spaces
- Uses of **term vectors**
  - Words can be compared, subtracted, found similarity of
  - **Word similarity** task: Are two words related?
  - **Word analogy** task: Given two related words and a third word, find the appropriate fourth word
- *Term-term context* matrix
  - **Context window**
  - Whether term  $i$  occurs within  $k$  number of terms upstream and downstream of term  $j$

# Term Co-occurrence and Context

- Global Vectors (GloVe)
- Global co-occurrence matrix  $X$
- $X_{ij}$  encodes how many times term  $i$  has appeared in the context of term  $j$
- Sum of a row  $X_i = \sum_{\forall j} X_{ij}$  denotes the *total* number of occurrences of term  $i$
- *Probability* that term  $j$  occurs in the context of term  $i$  is

$$P_{ij} = P(j|i) = \frac{X_{ij}}{X_i}$$

# Raw Probabilities and Ratio

- Raw counts  $X_{ij}$  or even raw probabilities  $P_{ij}$  may not be useful
- Depends a lot on the actual terms  $i$  and  $j$
- Also, asymmetric
- *Ratio* is a better indicator of relevance

	$P(k ice)$	$P(k steam)$	$\frac{P(k ice)}{P(k steam)}$
k = solid	1.9e-4	2.2e-5	8.636
k = gas	6.6e-5	7.8e-4	0.084
k = water	3.0e-3	2.2e-3	1.363
k = fashion	1.7e-5	1.8e-5	0.944

- “solid” is more relevant to “ice” than “steam”
- “gas” is more relevant to “steam” than “ice”
- “water” is relevant to both “ice” and “steam”
- “fashion” is irrelevant to both “ice” and “steam”

# Ratio of Probabilities

- Therefore, whether a term  $k$  is more relevant to term  $i$  than term  $j$  depends on the ratio  $P_{ik}/P_{jk}$
- Hence, for *context vector*  $\tilde{w}_k$  and *term vectors*  $w_i, w_j$

$$F(w_i, w_j, \tilde{w}_k) = \frac{P_{ik}}{P_{jk}}$$

- Vectors are of some dimensionality  $d$
- Since ratio of probabilities is scalar

$$F((w_i - w_j)^T \tilde{w}_k) = \frac{P_{ik}}{P_{jk}}$$

- Replacing probabilities by the same functional form

$$F((w_i - w_j)^T \tilde{w}_k) = \frac{F(w_i^T \tilde{w}_k)}{F(w_j^T \tilde{w}_k)}$$

# Symmetry

- Functional form solution is  $F = \exp$
- Since  $F(w_i^T \tilde{w}_k) = P_{ik}$

$$\begin{aligned}w_i^T \tilde{w}_k &= \log(P_{ik}) = \log(X_{ik}) - \log(X_i) \\w_i^T \tilde{w}_k + b_i &= \log(X_{ik})\end{aligned}$$

- Bias  $b_i$  encapsulates count of term  $i$
- Symmetric: both terms  $i$  and  $k$  should be in the context of each other

$$w_i^T \tilde{w}_k + b_i + \tilde{b}_k = \log(X_{ik})$$

- To avoid zero count problems

$$w_i^T \tilde{w}_k + b_i + \tilde{b}_k = \log(1 + X_{ik})$$

# Objective Function

- *Objective function* is weighted with (a function of)  $X_{ij}$

$$\arg \min_{w_i, w_j, \dots} \sum_{i,j=1}^V f(X_{ij}) \left( w_i^T \tilde{w}_j + b_i + \tilde{b}_j - \log(1 + X_{ij}) \right)^2$$

- Properties of weighting function  $f(X_{ij})$ 
  - $f(0) \rightarrow 0$
  - $f(X_{ij})$  is non-decreasing
  - $f(X_{ij})$  should not increase heavily for large values of  $x$
- Following function is used

$$f(X_{ij}) = \begin{cases} \left( \frac{X_{ij}}{X_{ij}^{max}} \right)^\alpha & \text{if } X_{ij} < X_{ij}^{max} \\ 1 & \text{otherwise} \end{cases}$$

- $\alpha = 3/4$
- $X_{ij}^{max} = 100$


# Discussion

- Context window of size  $\pm 5$
- Dimensionality of 100 or 300
- Captures *local context* and *global pairwise statistics*




# Perceptron

- A **perceptron** is a simple binary *linear* classifier
- Input attributes  $x_1, \dots, x_n$  are weighted and summed
- A bias  $b$  is added as well
- Final class ( $y = \pm 1$ ) is *sign* of the output
- The *sign* node is called the **activation function** or **link/decision/transfer function**
- Decision boundary is  $\vec{w} \cdot \vec{x} + b = \sum_{i=1}^n w_i x_i + b$
- Therefore, **sign** of  $\vec{w} \cdot \vec{x} + b$  predicts the class



- Why is bias needed?
- Otherwise, hyperplane passes through origin
- Simple trick to model input uniformly: include 1 as  $x_0$  of data
- Decision boundary becomes  $w_0x_0 + \vec{w} \cdot \vec{x} = \sum_{i=0}^n w_i x_i$
- Weight on  $x_0 = 1$  becomes the constant term, i.e.,  $w_0 = b$



# Examples of Perceptrons

- Different boolean functions
- AND (of  $x_1$  and  $x_2$ )
  - $w_1 = w_2 = 1, w_0 = -1.5$
- OR (of  $x_1$  and  $x_2$ )
  - $w_1 = w_2 = 1, w_0 = -0.5$
- NOT (of  $x_1$ )
  - $w_1 = -1, w_0 = 0.5$
- XOR (of  $x_1$  and  $x_2$ )
  - Cannot be done as the two classes are not linearly separable

# Learning a Perceptron

- What is new in a linear classifier?
- Training a perceptron, i.e., learning the weights  $w$
- **Perceptron learning rule** or **perceptron training rule**


$$w_i = w_i - \eta(\hat{y}_i - y_i)x_i$$

where  $\hat{y}_i$  is the predicted value and  $\eta$  is the *learning rate*


- If  $y_i = \hat{y}_i$ , there is no change in weight
- If  $y_i = +1$  and  $\hat{y}_i = -1$ , i.e.,  $\hat{y}_i - y_i < 0$ 
  - Weights of positive  $x_i$  are increased and those of negative  $x_i$  are decreased thereby pushing  $\hat{y}_i$  towards positive
- If  $y_i = -1$  and  $\hat{y}_i = +1$ , i.e.,  $\hat{y}_i - y_i > 0$ 
  - Weights of positive  $x_i$  are decreased and those of negative  $x_i$  are increased thereby pushing  $\hat{y}_i$  towards negative
- If the data *is* linearly separable, a perceptron *will* learn it, i.e., it will converge to the global optimum; otherwise, it may oscillate
- The learning rates  $\eta$  may be modified to give more importance to recent examples
- Weights can be learned through *gradient descent* method as well

# Combination of Perceptrons

- A single perceptron can learn only a single hyperplane
- If the data can be separated using two or more hyperplanes, a *combination* of perceptrons can learn it
- Example: XOR




- Each hyperplane can be modeled by a perceptron and the outputs can be combined using OR
  - *Multiple layers*



# Artificial Neural Networks

- **Artificial neural networks (ANNs)** are modeled on the human brain
- Nodes have connections ala neurons in human nervous system
- Nodes are also called **neurodes**
- Three types of nodes: input layer, *hidden* layer, output layer



- ANN with one hidden layer is considered as *two-layered*
  - Input layer is not counted
- Can have multiple hidden layers
- Learning through ANNs is also called **connectionist learning**

# Connections



- ANNs are of many types depending on the connections
- The most common are **multilayer feed-forward networks**
  - Connections are directed from one layer *only* to the next
  - There are *no* back edges or same-layer connections
- In **recurrent networks**, there can be back edges or same-layer connections
- **Fully connected**, i.e., each node in one layer is connected to every node in the next layer
- Training is learning the weights on each of these edges
- Akin to layers of perceptrons
- Activation functions in the nodes are *not* linear
  - Combination of linear functions can only learn linear separators
- Activation function is **sigmoid (logistic)** or **hyperbolic tangent**

# Sigmoid and Hyperbolic Tangent Functions

- If input of a node is  $x$ , then output  $y$  is

$$y = \sigma(x) = \frac{1}{1 + e^{-x}}$$

$$y = \tanh(x) = \frac{e^x - e^{-x}}{e^x + e^{-x}}$$



- Approximates the step (or sign) function
- Continuous and differentiable
- Output constrained to  $(0, 1)$  or  $(-1, +1)$
- Scaled versions of each other
- Also called **squashing functions**



# Nodes and Weights

- Output of a node is the sigmoid of the weighted sum of its inputs
- Inputs, in turn, are outputs of previous layers
- Inputs are normalized to  $(0, 1)$
- Outputs are already constrained to  $(0, 1)$
- Weight from a node  $i$  to node  $j$  is  $w_{ij}$
- Output from node  $i$  is  $O_i$
- Input to node  $j$  is  $I_j$

$$I_j = \sum_{\forall i} w_{ij} O_i$$


- Output from node  $j$  is  $O_j$

$$O_j = \sigma(I_j) = \frac{1}{1 + e^{-I_j}}$$

# Training an ANN

- Training an ANN requires
  - Designing the topology: how many layers and many nodes in each layer
  - Learning the weights of the connections
- Designing the topology requires either extensive domain knowledge or simply try-and-test
- The final outputs are some non-linear functions of inputs
- Hence, can be trained using gradient descent
- However, it is too complex and slow
- Weights are updated through the **backpropagation** algorithm

# Backpropagation



- Main idea

- Start with arbitrary weights
- Propagate forward the values
- Propagate backward the errors
- Update the weights using gradient descent

# Backpropagation

- Output, using sigmoid function  $\sigma(u) = \frac{1}{1+e^{-u}}$ , is

$$y = \sigma(u) = \sigma\left(\sum_{\forall x_i} w_i \cdot x_i\right)$$

- Derivatives of activation functions

$$\text{Sigmoid: } \frac{d\sigma(u)}{du} = \sigma(u)\sigma(-u) = \sigma(u)(1 - \sigma(u))$$

$$\text{Tanh: } \frac{d(\tanh(u))}{d(u)} = 1 - \tanh^2(u)$$

$$\text{Linear: } \frac{d(\text{ReLU}(u))}{d(u)} = \begin{cases} 0 & \text{when } u < 0 \\ 1 & \text{when } u \geq 0 \end{cases}$$

# Updating Weights

- If true output is  $t$ , error is squared error

$$E = \frac{1}{2}(t - y)^2$$

- *Stochastic gradient descent*
- Error function with respect to a single weight  $w_i$


$$\frac{\partial E}{\partial w_i} = \frac{\partial E}{\partial y} \cdot \frac{\partial y}{\partial u} \cdot \frac{\partial u}{\partial w_i} = [(y - t)] \cdot [y(1 - y)] \cdot [x_i]$$

- Therefore, update equation is

$$w_i^{(new)} = w_i^{(old)} - \eta \cdot (y - t) \cdot y(1 - y) \cdot x_i$$

- **Learning rate** or *momentum*  $\eta$  controls the speed of training
- Can be continued till error is below a threshold

# Backpropagation for Multiple Outputs



$$h_i = \sigma(u_i) = \sigma\left(\sum_{k=1}^K w_{ki} \cdot x_{ki}\right)$$

$$y_j = \sigma(u'_j) = \sigma\left(\sum_{i=1}^N w'_{ij} \cdot x_{ij}\right)$$

$$E(\vec{x}, \vec{t}, W, W') = \frac{1}{2} \sum_{j=1}^M (y_j - t_j)^2$$

# Representational Power of an ANN

- Boolean functions
  - Can approximate *any* boolean function with one layer of hidden nodes
  - Number of hidden nodes may be equal to exponential factor of number of boolean variables
  - Output layer wired through AND or OR
- Continuous functions
  - Can approximate *any* continuous function with one layer of hidden nodes up to any arbitrary error factor
  - Due to properties of sigmoid/tanh functions
  - Number of hidden nodes may be large
- Arbitrary functions
  - Can approximate *any* arbitrary function with *two* layers of hidden nodes up to any arbitrary error factor
  - Due to properties of sigmoid/tanh functions
  - Number of hidden nodes may be large

- Determining number of nodes and layers is problematic
- **Universal approximators**
  - Sigmoid and hyperbolic tangent functions
- Gradient descent converges to local minimum only
- Generally, requires large training data
- Very slow to train
- Can be easily parallelized
- Notoriously non-explainable



# Term Embedding


- Aim is to learn an embedding vector for each term that can *predict another term* in its context
- Thus, modeled as a *classification* problem
- ANNs used: **Word2Vec** model
- If vocabulary size is  $V$ , this is a  $V$ -class classification problem
- Input is a term *vector*
- Dimensionality is  $V$
- **One-hot vector**
  - Only the specific term is 1, rest are all 0
- Output layer consists of  $V$  nodes
- Only one output vector of dimensionality  $V$
- **Softmax** used for classification
- Probability of a particular dimension  $i$  with value  $f_i$  is

$$p(o_i) = \frac{\exp(f_i)}{\sum_{\forall i} \exp(f_i)}$$

- Choose the word  $w$  whose probability  $p(w)$  is the *highest*

# One-Word Context

- One word per context is used to predict one target word
  - Subhas : ? Bose



- Only one hidden layer of size  $N$
- Fully connected feed-forward architecture
- Number of weights is  $V \cdot N + N \cdot V$
- Given a one-hot encoded vector  $\vec{x}_I$  for input  $w_I$ , the output is  $\vec{y}_O$  predicting the word  $w_O$

# Layers

- Weight matrix  $W$  of size  $V \cdot N$  from input to hidden layer
- Row  $k$  represents vector for word  $k$ :  $v_k^T$
- For input vector  $x$  of word  $k$

$$h = W^T x = W^T I_{(k), \cdot} = v_{w_I}^T$$

- $v_{w_I}$  is the vector representation of input word  $w_I$  in  $N$  dimensions
- Essentially, the hidden layer is a *linear* copy
- Different weight matrix  $W'$  of size  $N \cdot V$  from hidden to output layer
- Score for each word  $w_j$  is

$$u_j = v_{w_j}'^T h = v_{w_j}'^T v_{w_I}$$

- $v_{w_j}'$  is the  $j$ -th column of  $W'$

# Objective

- Using softmax, a multinomial distribution over all predicted words  $w_j$  given an input word  $w_I$  is defined
- *Log-linear* classification model
- Probability of target word being  $w_j$  given an input context word  $w_I$  is

$$p(w_j|w_I) = y_j = \frac{\exp(u_j)}{\sum_{j'=1}^V \exp(u_{j'})} = \frac{\exp(v'_{w_j} v_{w_I})}{\sum_{j'=1}^V \exp(v'_{w_{j'}} v_{w_I})}$$


- $v_w$  is the *input* vector of word  $w$
- $v'_w$  is the *output* vector of word  $w$
- Training objective is to *maximize* the probabilities

$$\max p(w_{j_*}|w_I) = \max y_{j_*} = \max \log y_{j_*} = u_{j_*} - \log \sum_{j'=1}^V \exp(u_{j'}) = -E$$

- Which is the word vector representation?
- *Hidden layer* of dimensionality  $N$
- Training is through backpropagation

# Continuous Bag-of-Words (CBOW) Model

- **CBOW**: *Multiple context words* to predict a *single target word*
  - Tendulkar, Dravid, Laxman : ? Ganguly
  - Sachin, Rahul : ? DevVarman



- Bag-of-words in a local context window that continuously changes

- Instead of just a copy, hidden node is average of  $C$  context words


$$\vec{h} = \frac{1}{C} W^T (x_1 + x_2 + \dots + x_C) = \frac{1}{C} (v_{w_1} + v_{w_2} + \dots + v_{w_C})^T$$

- Error

$$E = -\log p(w_{j*} | w_{l,1}, \dots, w_{l,C}) = -u_{j*} + \log \sum_{j'=1}^V \exp(u'_{j'})$$

# Skip-Gram Model

- **Skip-gram**: *Single context word* to predict *multiple target words*
  - DevVarman : ? , ? Sachin, Rahul



- Output layer is  $C$  target words of dimensionality  $V$

# Model

- Output layer weights are shared, i.e., they are the same matrix  $W'$
- Probability of  $c^{\text{th}}$  target word being  $w_{c,j}$  given an input context word  $w_l$  is

$$p(w_{c,j}|w_l) = y_{c,j} = \frac{\exp(u_{c,j})}{\sum_{j'=1}^V \exp(u_{j'})} = \frac{\exp(v_{w_j}'^T v_{w_l})}{\sum_{j'=1}^V \exp(v_{w_{j'}}'^T v_{w_l})}$$

- Error

$$E = -\log p(w_{j^*,1}, \dots, w_{j^*,c}|w_l) = -\log \prod_{c=1}^C \frac{\exp(u_{c,j^*_c})}{\sum_{j'=1}^V \exp(u_{j'})}$$



# Word2Vec Discussion

- Context window size used is 4 words up and down
- Note that,  $k$  is chosen randomly within 4
  - This is why it is called “continuous” bag-of-words
- Number of context words,  $C$ , is within 5
- Rare words are discarded
  - This also effectively increases context window size
- Log-linear model
  - This makes training faster
- Input word and context word are treated differently
  - No good reason except mathematical convenience
- Skip-gram gives better semantic results
- *Corpus* is more important than method
- Why does just context?
  - No definitive or satisfying answer

# Morphologically Richer Languages

- Morphologically richer languages have the same internal root or structure for a large number of words
- Further, it requires a prohibitively large corpus to train for all such variants of a word
- Examples: Indian languages, Finnish, Turkish, and even European languages such as French, Spanish, German, Russian
- In Indian languages, compound words are a problem too
- Simple corpus-based word vector embeddings may not work
- Byte-pair encodings, etc. alleviate the problem
- **FastText** uses **character n-grams**
- Originally called **SISG (Subword Information Skip Gram)**

- A word vector representation is *sum* of character n-gram vectors and the word itself
- Two limits for n-gram length: minimum and maximum
  - Generally, 3 and 6
- Example word: “India”
- If limits are 2 and 3, then embeddings are of “In”, “Ind”, “ndi”, “dia”, “ia”, and “India”
- If  $G_w$  is the set of n-grams for a word  $w$ , then

$$s(w, c) = \sum_{\forall g \in G_w} z_g^T v_c$$

- FastText can handle *out-of-vocabulary* words
  - Glove and Word2Vec cannot!
- Uses as many common n-grams as it can from the new word